

[Get started](#)[Open in app](#)[Follow](#)

600K Followers



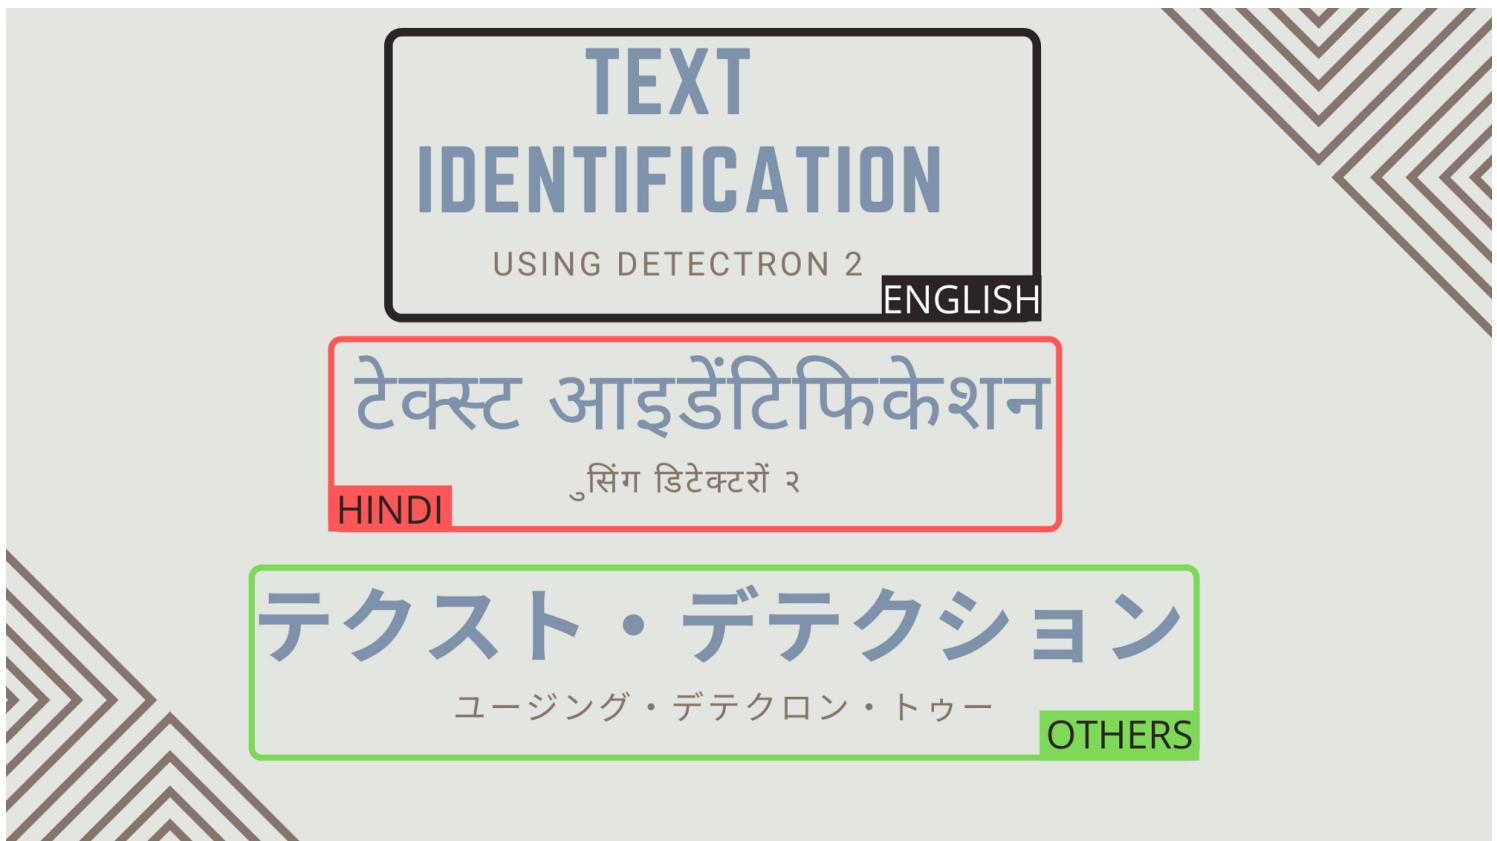
You have **2** free member-only stories left this month. [Sign up for Medium and get an extra one](#)

Object Detection in 6 steps using Detectron2

Let's look at how to use FAIR's (Facebook AI Research) Detectron 2 for Instance Detection on a custom dataset which involves text identification.



Aakarsh Yelisetty Aug 2, 2020 · 7 min read ★



Designed using Canva

Have you ever tried training an object detection model using a custom dataset of your own choice from scratch?

If yes, you'd know how tedious the process would be. We need to start with building a model using a Feature Pyramid Network combined with a Region Proposal Network if we opt for region proposal based methods such as Faster R-CNN or we can also use one-shot detector algorithms like SSD and YOLO.

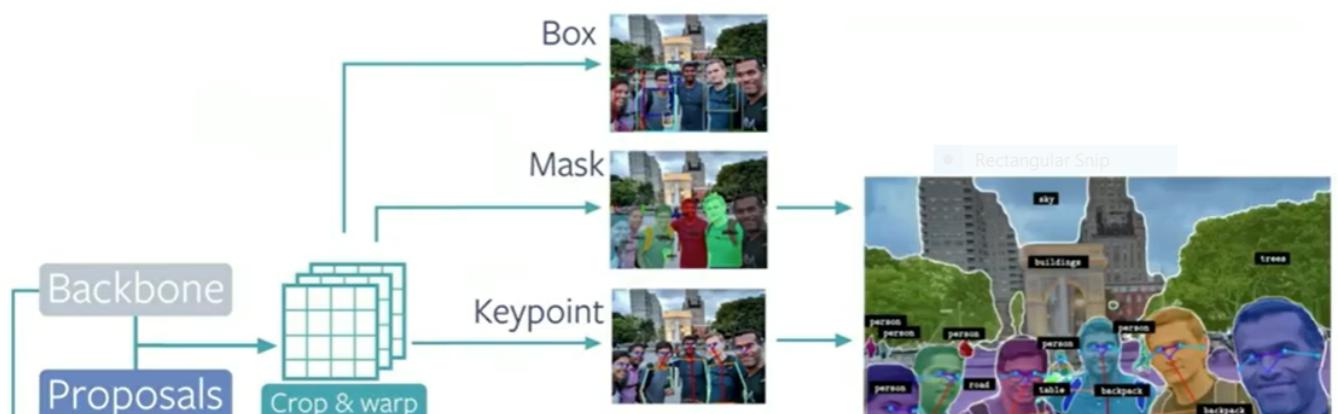
Either of them is a bit complicated to work with if we want to implement it from scratch. We need a framework where we can use state-of-the-art models such as Fast, Faster, and Mask R-CNNs with ease. Nevertheless, it is important to try building a model at least once from scratch to understand the math behind it.

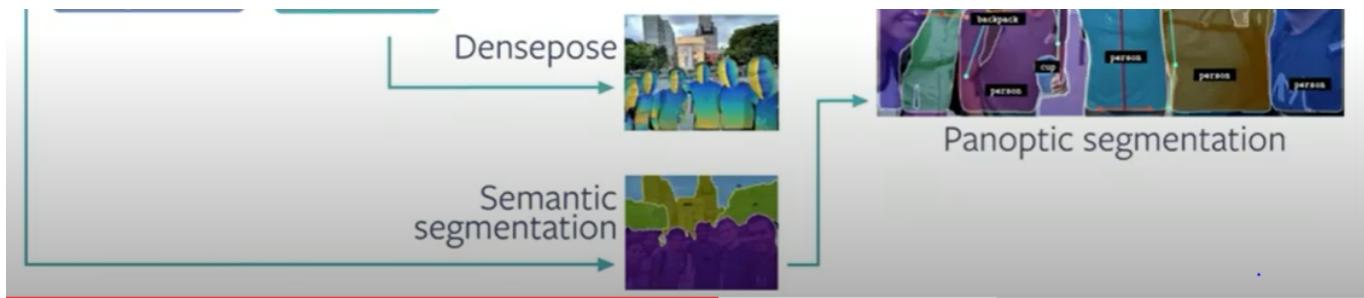
Detectron 2 comes to the rescue if we want to train an object detection model in a snap with a custom dataset. All the models present in the [model zoo](#) of the Detectron 2 library are pre-trained on COCO Dataset. We just need to fine-tune our custom dataset on the pre-trained model.

Detectron 2 is a complete rewrite of the first Detectron which was released in the year 2018. The predecessor was written on Caffe2, a deep learning framework that is also backed by Facebook. Both the Caffe2 and Detectron are now deprecated. Caffe2 is now a part of PyTorch and the successor, Detectron 2 is completely written on PyTorch.

Detectron2 is meant to advance machine learning by offering speedy training and addressing the issues companies face when making the step from research to production.

These are the various types of Object Detection models that the Detectron 2 offers.





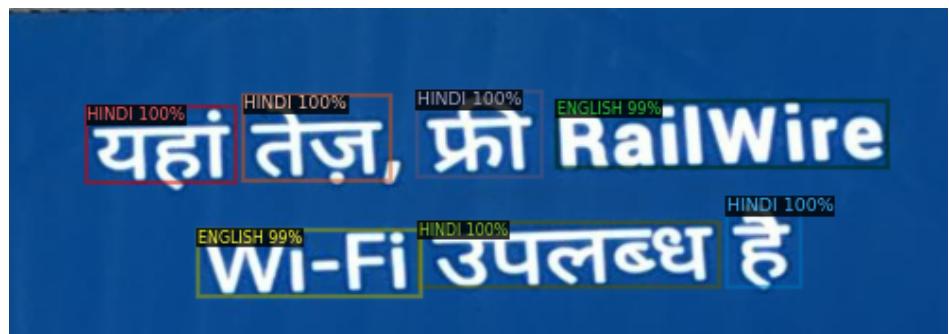
<https://research.fb.com/wp-content/uploads/2019/12/4.-detectron2.pdf>

Let's dive into **Instance Detection** directly.

Instance Detection refers to the classification and localization of an object with a bounding box around it. In this article, We are going to deal with identifying the language of text from images using the Faster RCNN model from the Detectron 2's model zoo.

Note that we are going to limit our languages by 2.

We identify **Hindi** and **English** Text and we include a class labeled **Others** for other languages.



Final Results from Colab

We'll implement a model where the output would be in this way.

Let's get started!

Using Detectron 2, Object Detection can be performed on any custom dataset using seven steps. All the steps are readily available in this [Google Colab Notebook](#) and you can run it straight away!

Using Google Colab for this would be an easy task as we can use a GPU for faster training.

Step 1: Installing Detectron 2

Kickstart with installing a few dependencies such as Torch Vision and COCO API and check whether **CUDA** is available. CUDA helps in keeping track of the currently selected GPU. And then install Detectron2.

```
# install dependencies:  
!pip install -U torch==1.5 torchvision==0.6 -f  
https://download.pytorch.org/whl/cu101/torch\_stable.html  
!pip install cython pyyaml==5.1  
!pip install -U  
'git+https://github.com/cocodataset/cocoapi.git#subdirectory=PythonAPI'  
  
import torch, torchvision  
print(torch.__version__, torch.cuda.is_available())  
!gcc --version  
  
# install detectron2:  
!pip install detectron2==0.1.3 -f  
https://dl.fbaipublicfiles.com/detectron2/wheels/cu101/torch1.5/index.html
```

Step 2: Prepare and Register the Dataset

Import a few necessary packages.

```
1 # You may need to restart your runtime prior to this, to let your installation take effect  
2 import detectron2  
3 from detectron2.utils.logger import setup_logger  
4 setup_logger()  
5  
6 # import some common libraries  
7 import numpy as np  
8 import cv2
```

```

1   import cv2
2
3   import random
4
5   from google.colab.patches import cv2_imshow
6
7
8   # import some common detectron2 utilities
9   from detectron2 import model_zoo
10  from detectron2.engine import DefaultPredictor
11  from detectron2.config import get_cfg
12  from detectron2.utils.visualizer import Visualizer
13  from detectron2.data import MetadataCatalog

```

imports.py hosted with ❤ by GitHub

[view raw](#)

Datasets that have builtin support in detectron2 are listed in [builtin datasets](#). If you want to use a custom dataset while also reusing detectron2's data loaders, you will need to **Register** your dataset (i.e., tell detectron2 how to obtain your dataset).

We use our Text Detection Dataset which has three classes:

1. English
2. Hindi
3. Others

We'll train a text detection model from an existing model pre-trained on the COCO dataset, available in detectron2's model zoo.

If you're interested to know the conversion from the raw format of the dataset to the format that Detectron 2 accepts, checkout this [Colab Notebook](#).

There are certain formats in how data can be fed to a model such as a YOLO format, PASCAL VOC format, COCO format, etc. Detectron2 accepts the COCO Format of the dataset. COCO format of the dataset consists of a JSON file which includes all the details of an image such as size, annotations (i.e., bounding box coordinates), labels corresponding to it's bounding box, etc. For example,

```

1  {
2      "image_id": 0,
3      "file_name": "401.jpeg",
4      "height": 500

```

```

5      "width": 715,
6      "annotations": [
7      {
8          "bbox": [
9              67,
10             200,
11             187,
12             59
13         ],
14         "bbox_mode": 1,
15         "category_id": "0"
16     },
17     {
18         "bbox": [
19             279,
20             170,
21             417,
22             38
23         ],
24         "bbox_mode": 1,
25         "category_id": "0"
26     }
27 ]
28 }

```

Detectron 2's format of Dataset

This is how the JSON looks like for one image. There are different types of formats for the bounding box representation. It must be a member of `structures.BoxMode` for Detectron2. There are 5 such formats. But currently, it supports `BoxMode.XYXY_ABS`, `BoxMode.XYWH_ABS`. We use the second format. The (X, Y) represents one coordinate of the bounding box, and W, H represents Width and Height of that box. The `category_id` refers to the class the box belongs to.

Then, we need to register our dataset.

```

1 import json
2 from detectron2.structures import BoxMode
3 def get_board_dicts(imgdir):
4     json_file = imgdir + "/dataset.json" #Fetch the json file
5     with open(json_file) as f:
6         dataset_dicts = json.load(f)
7     for i in dataset_dicts:

```

```

8     filename = i["file_name"]
9     i["file_name"] = imgdir+"/"+filename
10    for j in i["annotations"]:
11        j["bbox_mode"] = BoxMode.XYWH_ABS #Setting the required Box Mode
12        j["category_id"] = int(j["category_id"])
13    return dataset_dicts
14 from detectron2.data import DatasetCatalog, MetadataCatalog
15 #Registering the Dataset
16 for d in ["train", "val"]:
17     DatasetCatalog.register("boardetect_" + d, lambda d=d: get_board_dicts("Text_Detecti
18     MetadataCatalog.get("boardetect_" + d).set(thing_classes=["HINDI","ENGLISH","OTHER"])
19 board_metadata = MetadataCatalog.get("boardetect_train")

```

get_board_dicts.py hosted with ❤ by GitHub

[view raw](#)

To verify the data loading is correct, let's visualize the annotations of randomly selected samples in the training set.

Step 3: Visualize the Training Set

We'll randomly pick 3 pictures from the train folder of our dataset and see how the bounding boxes look like.

```

1 #Visualizing the Train Dataset
2 dataset_dicts = get_board_dicts("Text_Detection_Dataset_COCO_Format/train")
3 #Randomly choosing 3 images from the Set
4 for d in random.sample(dataset_dicts, 3):
5     img = cv2.imread(d["file_name"])
6     visualizer = Visualizer(img[:, :, ::-1], metadata=board_metadata)
7     vis = visualizer.draw_dataset_dict(d)
8     cv2.imshow(vis.get_image()[:, :, ::-1])

```

visualize.py hosted with ❤ by GitHub

[view raw](#)

The output looks in this way,



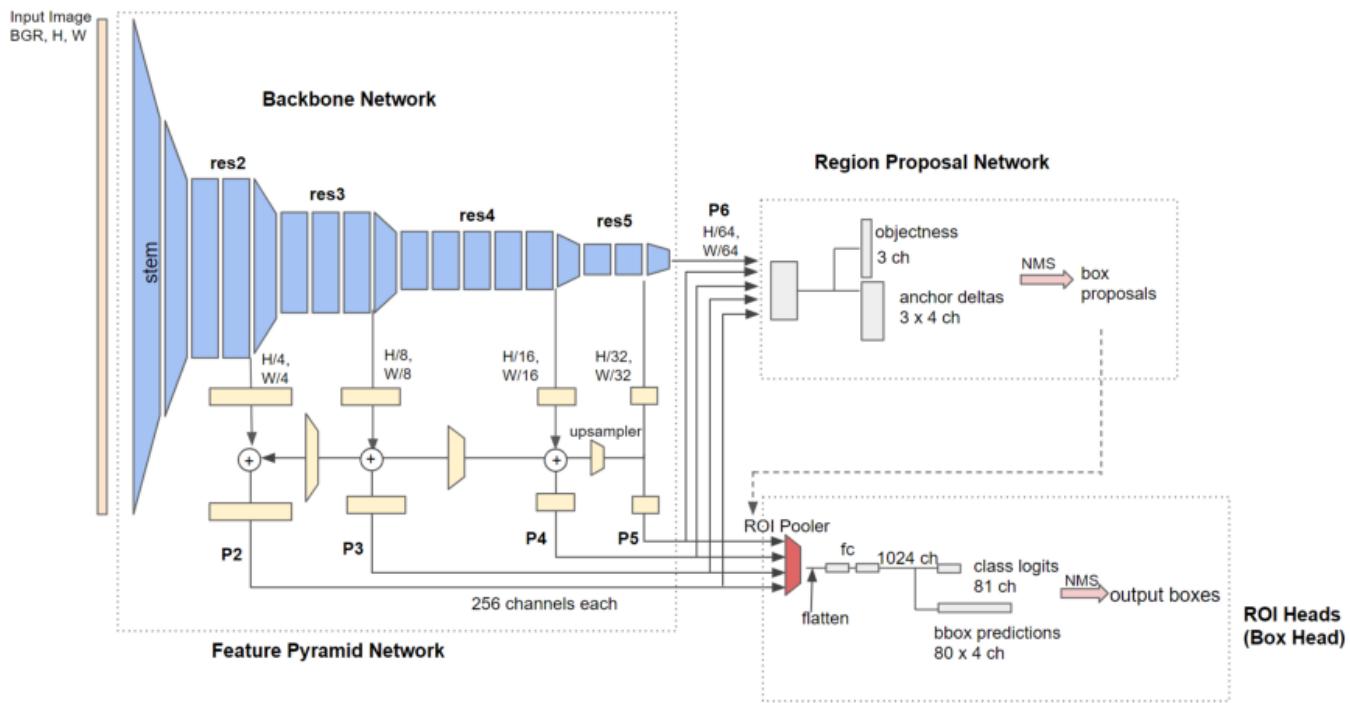
Results from Colab

Step 4: Training the Model

The big step. This is the step where we give configurations and set the model ready to get trained. Technically, we just fine-tune our model on the dataset as the model is

already pre-trained on COCO Dataset.

There are a ton of models available for object detection in the Detectron2's Model Zoo. Here, we use the **faster_rcnn_R_50_FPN_3x** model which looks in this way on a high level.



[Source](#)

There'd be a Backbone Network (Resnet in this case) which is used to extract features from the image followed by a Region Proposal Network for proposing region proposals and a Box Head for tightening the bounding box.

You can read more about how Faster R-CNN works in my [previous article](#).

Let's set the configuration for training.

```

1  from detectron2.engine import DefaultTrainer
2  from detectron2.config import get_cfg
3  import os
4  cfg = get_cfg()
5  cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/faster_rcnn_R_50_FPN_3x.yaml"))
6  #Passing the Train and Validation sets
7  cfg.DATASETS.TRAIN = ("boardetect_train",)
8  cfg.DATASETS.TEST = ("boardetect_val",)

```

```

9  # Number of data loading threads
10 cfg.DATALOADER.NUM_WORKERS = 4
11 cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/faster_rcnn_R_50_FPN_3x")
12 # Number of images per batch across all machines.
13 cfg.SOLVERIMS_PER_BATCH = 4
14 cfg.SOLVER.BASE_LR = 0.0125 # pick a good LearningRate
15 cfg.SOLVER.MAX_ITER = 1500 #No. of iterations
16 cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 256
17 cfg.MODEL.ROI_HEADS.NUM_CLASSES = 3 # No. of classes = [HINDI, ENGLISH, OTHER]
18 cfg.TEST.EVAL_PERIOD = 500 # No. of iterations after which the Validation Set is evaluated
19 os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
20 trainer = CocoTrainer(cfg)
21 trainer.resume_or_load(resume=False)
22 trainer.train()

```

set_config.py hosted with ❤ by GitHub

[view raw](#)

I wouldn't say this is the best configuration. Surely, the accuracy may get improved for other configs as well. After all, it depends on choosing the right hyperparameters.

```

cfg.MODEL.ROI_HEADS.NUM_CLASSES = 3 # No. of classes = [HINDI, ENGLISH, OTHER]
cfg.TEST.EVAL_PERIOD = 500
os.makedirs(cfg.OUTPUT_DIR, exist_ok=True)
trainer = CocoTrainer(cfg)
trainer.resume_or_load(resume=False)
trainer.train()

[08/01 12:28:15 d2.utils.events]: eta: 0:25:53 iter: 39 total_loss: 1.522 loss_cls: 0.602 loss_box_reg: 0.758 loss_rpn_cls: 0.055 loss_rpn_loc: 0.056 time: 1.0614 data_time
[08/01 12:28:36 d2.utils.events]: eta: 0:25:06 iter: 59 total_loss: 1.295 loss_cls: 0.395 loss_box_reg: 0.712 loss_rpn_cls: 0.044 loss_rpn_loc: 0.062 time: 1.0534 data_time
[08/01 12:28:57 d2.utils.events]: eta: 0:24:48 iter: 79 total_loss: 1.110 loss_cls: 0.366 loss_box_reg: 0.632 loss_rpn_cls: 0.026 loss_rpn_loc: 0.051 time: 1.0476 data_time
[08/01 12:29:18 d2.utils.events]: eta: 0:24:31 iter: 99 total_loss: 1.007 loss_cls: 0.351 loss_box_reg: 0.478 loss_rpn_cls: 0.041 loss_rpn_loc: 0.072 time: 1.0517 data_time
[08/01 12:29:39 d2.utils.events]: eta: 0:24:14 iter: 119 total_loss: 0.804 loss_cls: 0.251 loss_box_reg: 0.400 loss_rpn_cls: 0.029 loss_rpn_loc: 0.058 time: 1.0518 data_time
[08/01 12:30:00 d2.utils.events]: eta: 0:23:48 iter: 139 total_loss: 0.738 loss_cls: 0.230 loss_box_reg: 0.397 loss_rpn_cls: 0.030 loss_rpn_loc: 0.053 time: 1.0483 data_time
[08/01 12:30:21 d2.utils.events]: eta: 0:23:28 iter: 159 total_loss: 0.694 loss_cls: 0.229 loss_box_reg: 0.383 loss_rpn_cls: 0.027 loss_rpn_loc: 0.050 time: 1.0494 data_time
[08/01 12:30:41 d2.utils.events]: eta: 0:23:03 iter: 179 total_loss: 0.688 loss_cls: 0.224 loss_box_reg: 0.351 loss_rpn_cls: 0.024 loss_rpn_loc: 0.039 time: 1.0466 data_time
[08/01 12:31:03 d2.utils.events]: eta: 0:22:45 iter: 199 total_loss: 0.758 loss_cls: 0.247 loss_box_reg: 0.388 loss_rpn_cls: 0.022 loss_rpn_loc: 0.046 time: 1.0475 data_time
[08/01 12:31:23 d2.utils.events]: eta: 0:22:21 iter: 219 total_loss: 0.631 loss_cls: 0.221 loss_box_reg: 0.350 loss_rpn_cls: 0.015 loss_rpn_loc: 0.043 time: 1.0467 data_time
[08/01 12:31:44 d2.utils.events]: eta: 0:22:00 iter: 239 total_loss: 0.654 loss_cls: 0.200 loss_box_reg: 0.346 loss_rpn_cls: 0.017 loss_rpn_loc: 0.044 time: 1.0454 data_time
[08/01 12:32:05 d2.utils.events]: eta: 0:21:38 iter: 259 total_loss: 0.581 loss_cls: 0.194 loss_box_reg: 0.340 loss_rpn_cls: 0.017 loss_rpn_loc: 0.044 time: 1.0444 data_time
[08/01 12:32:25 d2.utils.events]: eta: 0:21:20 iter: 279 total_loss: 0.613 loss_cls: 0.184 loss_box_reg: 0.355 loss_rpn_cls: 0.016 loss_rpn_loc: 0.041 time: 1.0439 data_time
[08/01 12:32:46 d2.utils.events]: eta: 0:20:56 iter: 299 total_loss: 0.680 loss_cls: 0.217 loss_box_reg: 0.359 loss_rpn_cls: 0.025 loss_rpn_loc: 0.057 time: 1.0440 data_time
[08/01 12:33:08 d2.utils.events]: eta: 0:20:39 iter: 319 total_loss: 0.599 loss_cls: 0.197 loss_box_reg: 0.363 loss_rpn_cls: 0.015 loss_rpn_loc: 0.043 time: 1.0470 data_time
[08/01 12:33:29 d2.utils.events]: eta: 0:20:18 iter: 339 total_loss: 0.551 loss_cls: 0.172 loss_box_reg: 0.324 loss_rpn_cls: 0.010 loss_rpn_loc: 0.035 time: 1.0468 data_time
[08/01 12:33:50 d2.utils.events]: eta: 0:19:57 iter: 359 total_loss: 0.543 loss_cls: 0.150 loss_box_reg: 0.309 loss_rpn_cls: 0.010 loss_rpn_loc: 0.050 time: 1.0467 data_time
[08/01 12:34:12 d2.utils.events]: eta: 0:19:38 iter: 379 total_loss: 0.565 loss_cls: 0.163 loss_box_reg: 0.295 loss_rpn_cls: 0.024 loss_rpn_loc: 0.037 time: 1.0480 data_time
[08/01 12:34:31 d2.utils.events]: eta: 0:19:14 iter: 399 total_loss: 0.620 loss_cls: 0.212 loss_box_reg: 0.326 loss_rpn_cls: 0.016 loss_rpn_loc: 0.036 time: 1.0448 data_time
[08/01 12:34:53 d2.utils.events]: eta: 0:18:54 iter: 419 total_loss: 0.574 loss_cls: 0.167 loss_box_reg: 0.323 loss_rpn_cls: 0.017 loss_rpn_loc: 0.043 time: 1.0460 data_time
[08/01 12:35:14 d2.utils.events]: eta: 0:18:34 iter: 439 total_loss: 0.561 loss_cls: 0.160 loss_box_reg: 0.301 loss_rpn_cls: 0.013 loss_rpn_loc: 0.040 time: 1.0468 data_time
[08/01 12:35:34 d2.utils.events]: eta: 0:18:10 iter: 459 total_loss: 0.543 loss_cls: 0.171 loss_box_reg: 0.332 loss_rpn_cls: 0.010 loss_rpn_loc: 0.034 time: 1.0440 data_time
[08/01 12:35:54 d2.utils.events]: eta: 0:17:48 iter: 479 total_loss: 0.532 loss_cls: 0.172 loss_box_reg: 0.329 loss_rpn_cls: 0.009 loss_rpn_loc: 0.038 time: 1.0438 data_time

```

Training Process (Results from Colab)

Note that, here we also calculate the accuracy for every 500 iterations on the Validation Set.

Step 5: Inference using the Trained Model

It's time to infer the results by testing the model on the Validation Set.

An output folder gets saved in the local storage after successful completion of training in which the final weights are stored. You can save this folder for inferencing from this model in the future.

```

1  from detectron2.utils.visualizer import ColorMode
2
3  #Use the final weights generated after successful training for inference
4  cfg.MODEL.WEIGHTS = os.path.join(cfg.OUTPUT_DIR, "model_final.pth")
5
6  cfg.MODEL.ROI_HEADS.SCORE_THRESH_TEST = 0.8 # set the testing threshold for this model
7  #Pass the validation dataset
8  cfg.DATASETS.TEST = ("boardetect_val", )
9
10 predictor = DefaultPredictor(cfg)
11
12 dataset_dicts = get_board_dicts("Text_Detection_Dataset_COCO_Format/val")
13 for d in random.sample(dataset_dicts, 3):
14     im = cv2.imread(d["file_name"])
15     outputs = predictor(im)
16     v = Visualizer(im[:, :, ::-1],
17                     metadata=board_metadata,
18                     scale=0.8,
19                     instance_mode=ColorMode.IMAGE
20     )
21     v = v.draw_instance_predictions(outputs["instances"].to("cpu")) #Passing the predict
22     cv2.imshow(v.get_image()[:, :, ::-1])

```

[predictions.py](#) hosted with ❤ by GitHub

[view raw](#)

Results:



Step 6: Evaluation of the Trained Model

Usually, the model is evaluated following the COCO Standards of evaluation. Mean Average Precision (mAP) is used to evaluate the performance of the model. Here's an [article](#) that precisely gives an idea on the mAP.

```

1 #import the COCO Evaluator to use the COCO Metrics
2 from detectron2.evaluation import COCOEvaluator, inference_on_dataset
3 from detectron2.data import build_detection_test_loader
4
5 #Call the COCO Evaluator function and pass the Validation Dataset
6 evaluator = COCOEvaluator("boardetect_val", cfg, False, output_dir="/output/")
7 val_loader = build_detection_test_loader(cfg, "boardetect_val")
8
9 #Use the created predicted model in the previous step
10 inference_on_dataset(predictor.model, val_loader, evaluator)

```

[evaluation.py](#) hosted with ❤️ by GitHub

[view raw](#)

```

[07/28 12:59:29 d2.data.common]: Serializing 27 elements to byte tensors and concatenating them all ...
[07/28 12:59:29 d2.data.common]: Serialized dataset takes 0.01 MiB
[07/28 12:59:29 d2.evaluation.evaluator]: Start inference on 27 images
[07/28 12:59:33 d2.evaluation.evaluator]: Inference done 11/27. 0.3363 s / img. ETA=0:00:05
[07/28 12:59:39 d2.evaluation.evaluator]: Inference done 27/27. 0.3246 s / img. ETA=0:00:00
[07/28 12:59:39 d2.evaluation.evaluator]: Total inference time: 0:00:07.273351 (0.330607 s / img per device, on 1 devices)
[07/28 12:59:39 d2.evaluation.evaluator]: Total inference pure compute time: 0:00:07 (0.324569 s / img per device, on 1 devices)
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Preparing results for COCO format ...
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Saving results to /output/coco_instances_results.json
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Evaluating predictions ...
Loading and preparing results...
DONE (t=0.00s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *bbox*
DONE (t=0.07s).
Accumulating evaluation results...
DONE (t=0.03s).
Average Precision (AP) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.486
Average Precision (AP) @[ IoU=0.50    | area=   all | maxDets=100 ] = 0.794
Average Precision (AP) @[ IoU=0.75    | area=   all | maxDets=100 ] = 0.588
Average Precision (AP) @[ IoU=0.50:0.95 | area= small  | maxDets=100 ] = 0.416
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.541
Average Precision (AP) @[ IoU=0.50:0.95 | area= large  | maxDets=100 ] = 0.454
Average Recall (AR) @[ IoU=0.50:0.95 | area=   all | maxDets= 1 ] = 0.195
Average Recall (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=10 ] = 0.536
Average Recall (AR) @[ IoU=0.50:0.95 | area=   all | maxDets=100 ] = 0.538
Average Recall (AR) @[ IoU=0.50:0.95 | area= small  | maxDets=100 ] = 0.498
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.610
Average Recall (AR) @[ IoU=0.50:0.95 | area= large  | maxDets=100 ] = 0.519
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP      | AP50   | AP75   | APs    | APm   | APl   |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| 48.614 | 79.410 | 58.818 | 41.596 | 54.075 | 45.368 |
[07/28 12:59:39 d2.evaluation.coco_evaluation]: Per-category bbox AP:
| category | AP     | category | AP     | category | AP     |
|:-----:|:-----:|:-----:|:-----:|:-----:|:-----:|
| HINDI   | 57.467 | ENGLISH | 45.207 | OTHER    | 43.168 |
OrderedDict([('bbox',
{'AP': 48.614130895033696,
'AP-ENGLISH': 45.20707317037124,

```

Evaluation Metrics (Results from Colab)

We get an accuracy of around 79.4% for an IoU of 0.5 which is not that bad. It can be increased certainly by tweaking the parameters a bit and increasing the number of iterations. But keep an eye on training extensively as the model might overfit the training set.

Go through this [Colab Notebook](#) if you need inference from a saved model.

Conclusion

In this article, I emphasized on the procedure of objection detection using a custom dataset using detectron 2 rather than focusing on gaining higher accuracy.

Although this seems to be a pretty straight forward process, there's a lot to explore in Detectron 2's library. We have a good amount of optimization parameters that can be further tweaked to attain higher accuracy, which completely depends on one's custom dataset.

Hope you've learned something new today.

You can download the notebook from my [Github repository](#) and try running it on Google Colab or Jupyter Notebooks.

You could read more about Object Detection and the legacy R-CNN's in my [previous article](#).

If you'd like to get in touch, connect with me on [LinkedIn](#).

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

Get this newsletter

[Artificial Intelligence](#)[Computer Vision](#)[Object Detection](#)[Deep Learning](#)[Towards Data Science](#)[About](#) [Write](#) [Help](#) [Legal](#)[Get the Medium app](#)